

RETINAL DISEASE CLASSIFICATION

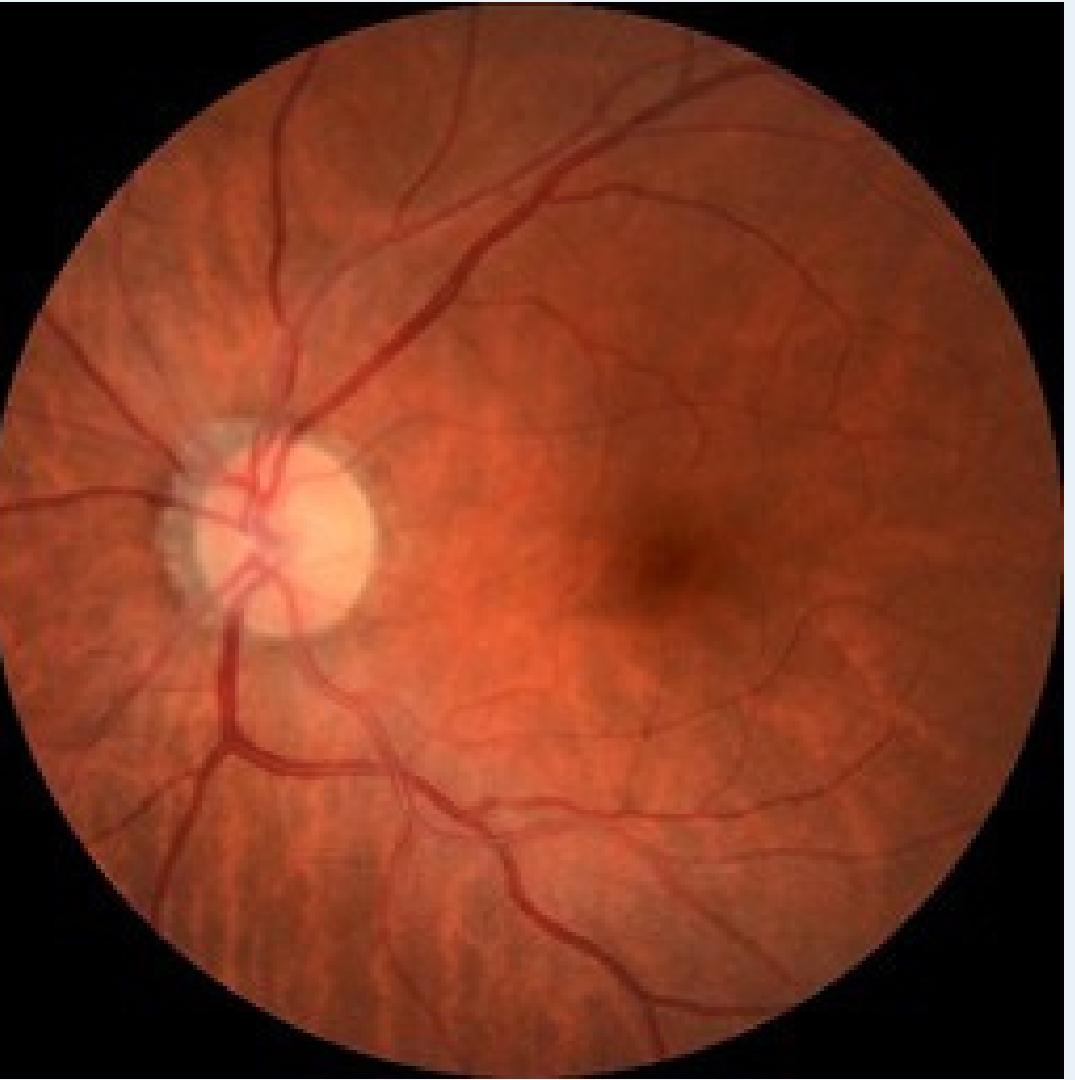
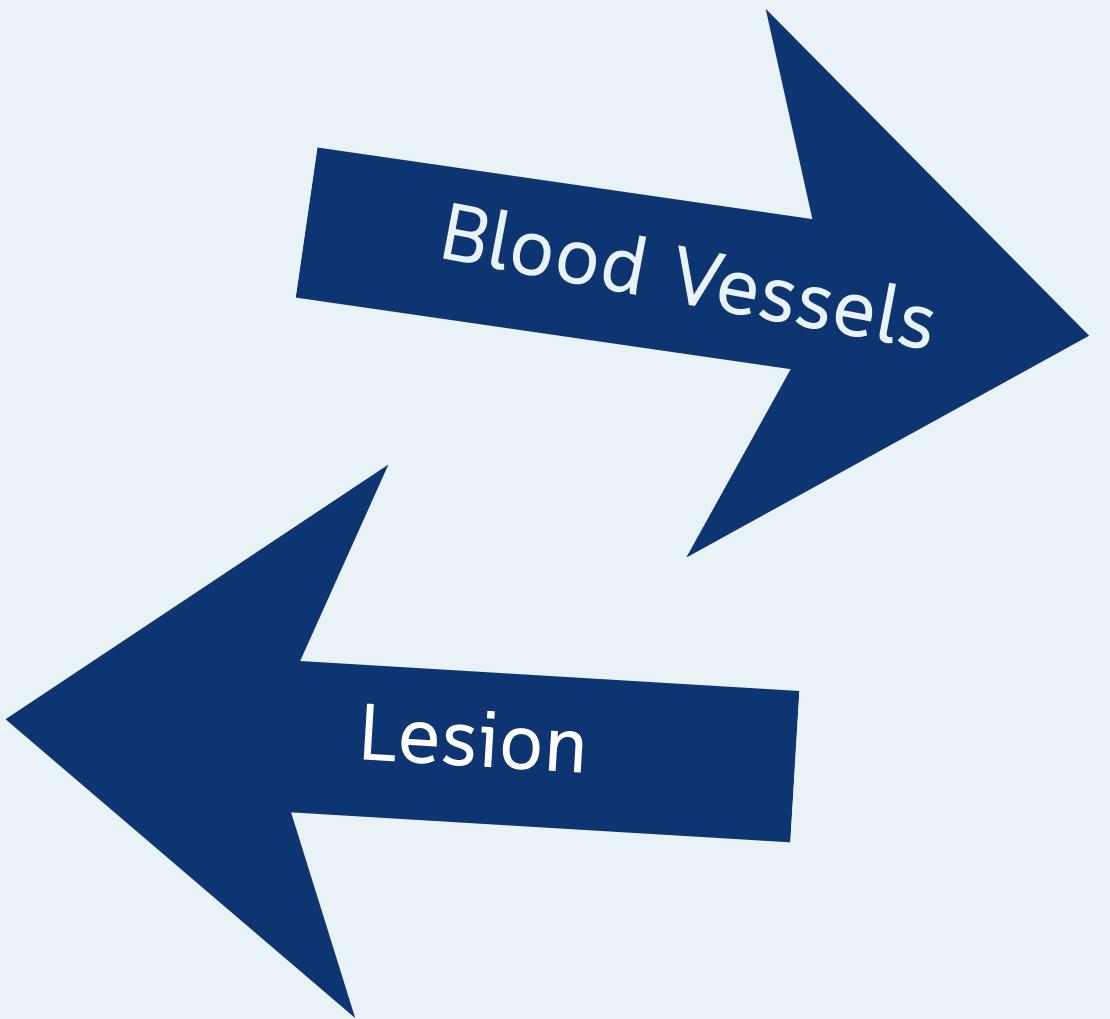
GROUP 20

Nguyen Quang
Hoai Thuong
Viet Hoang

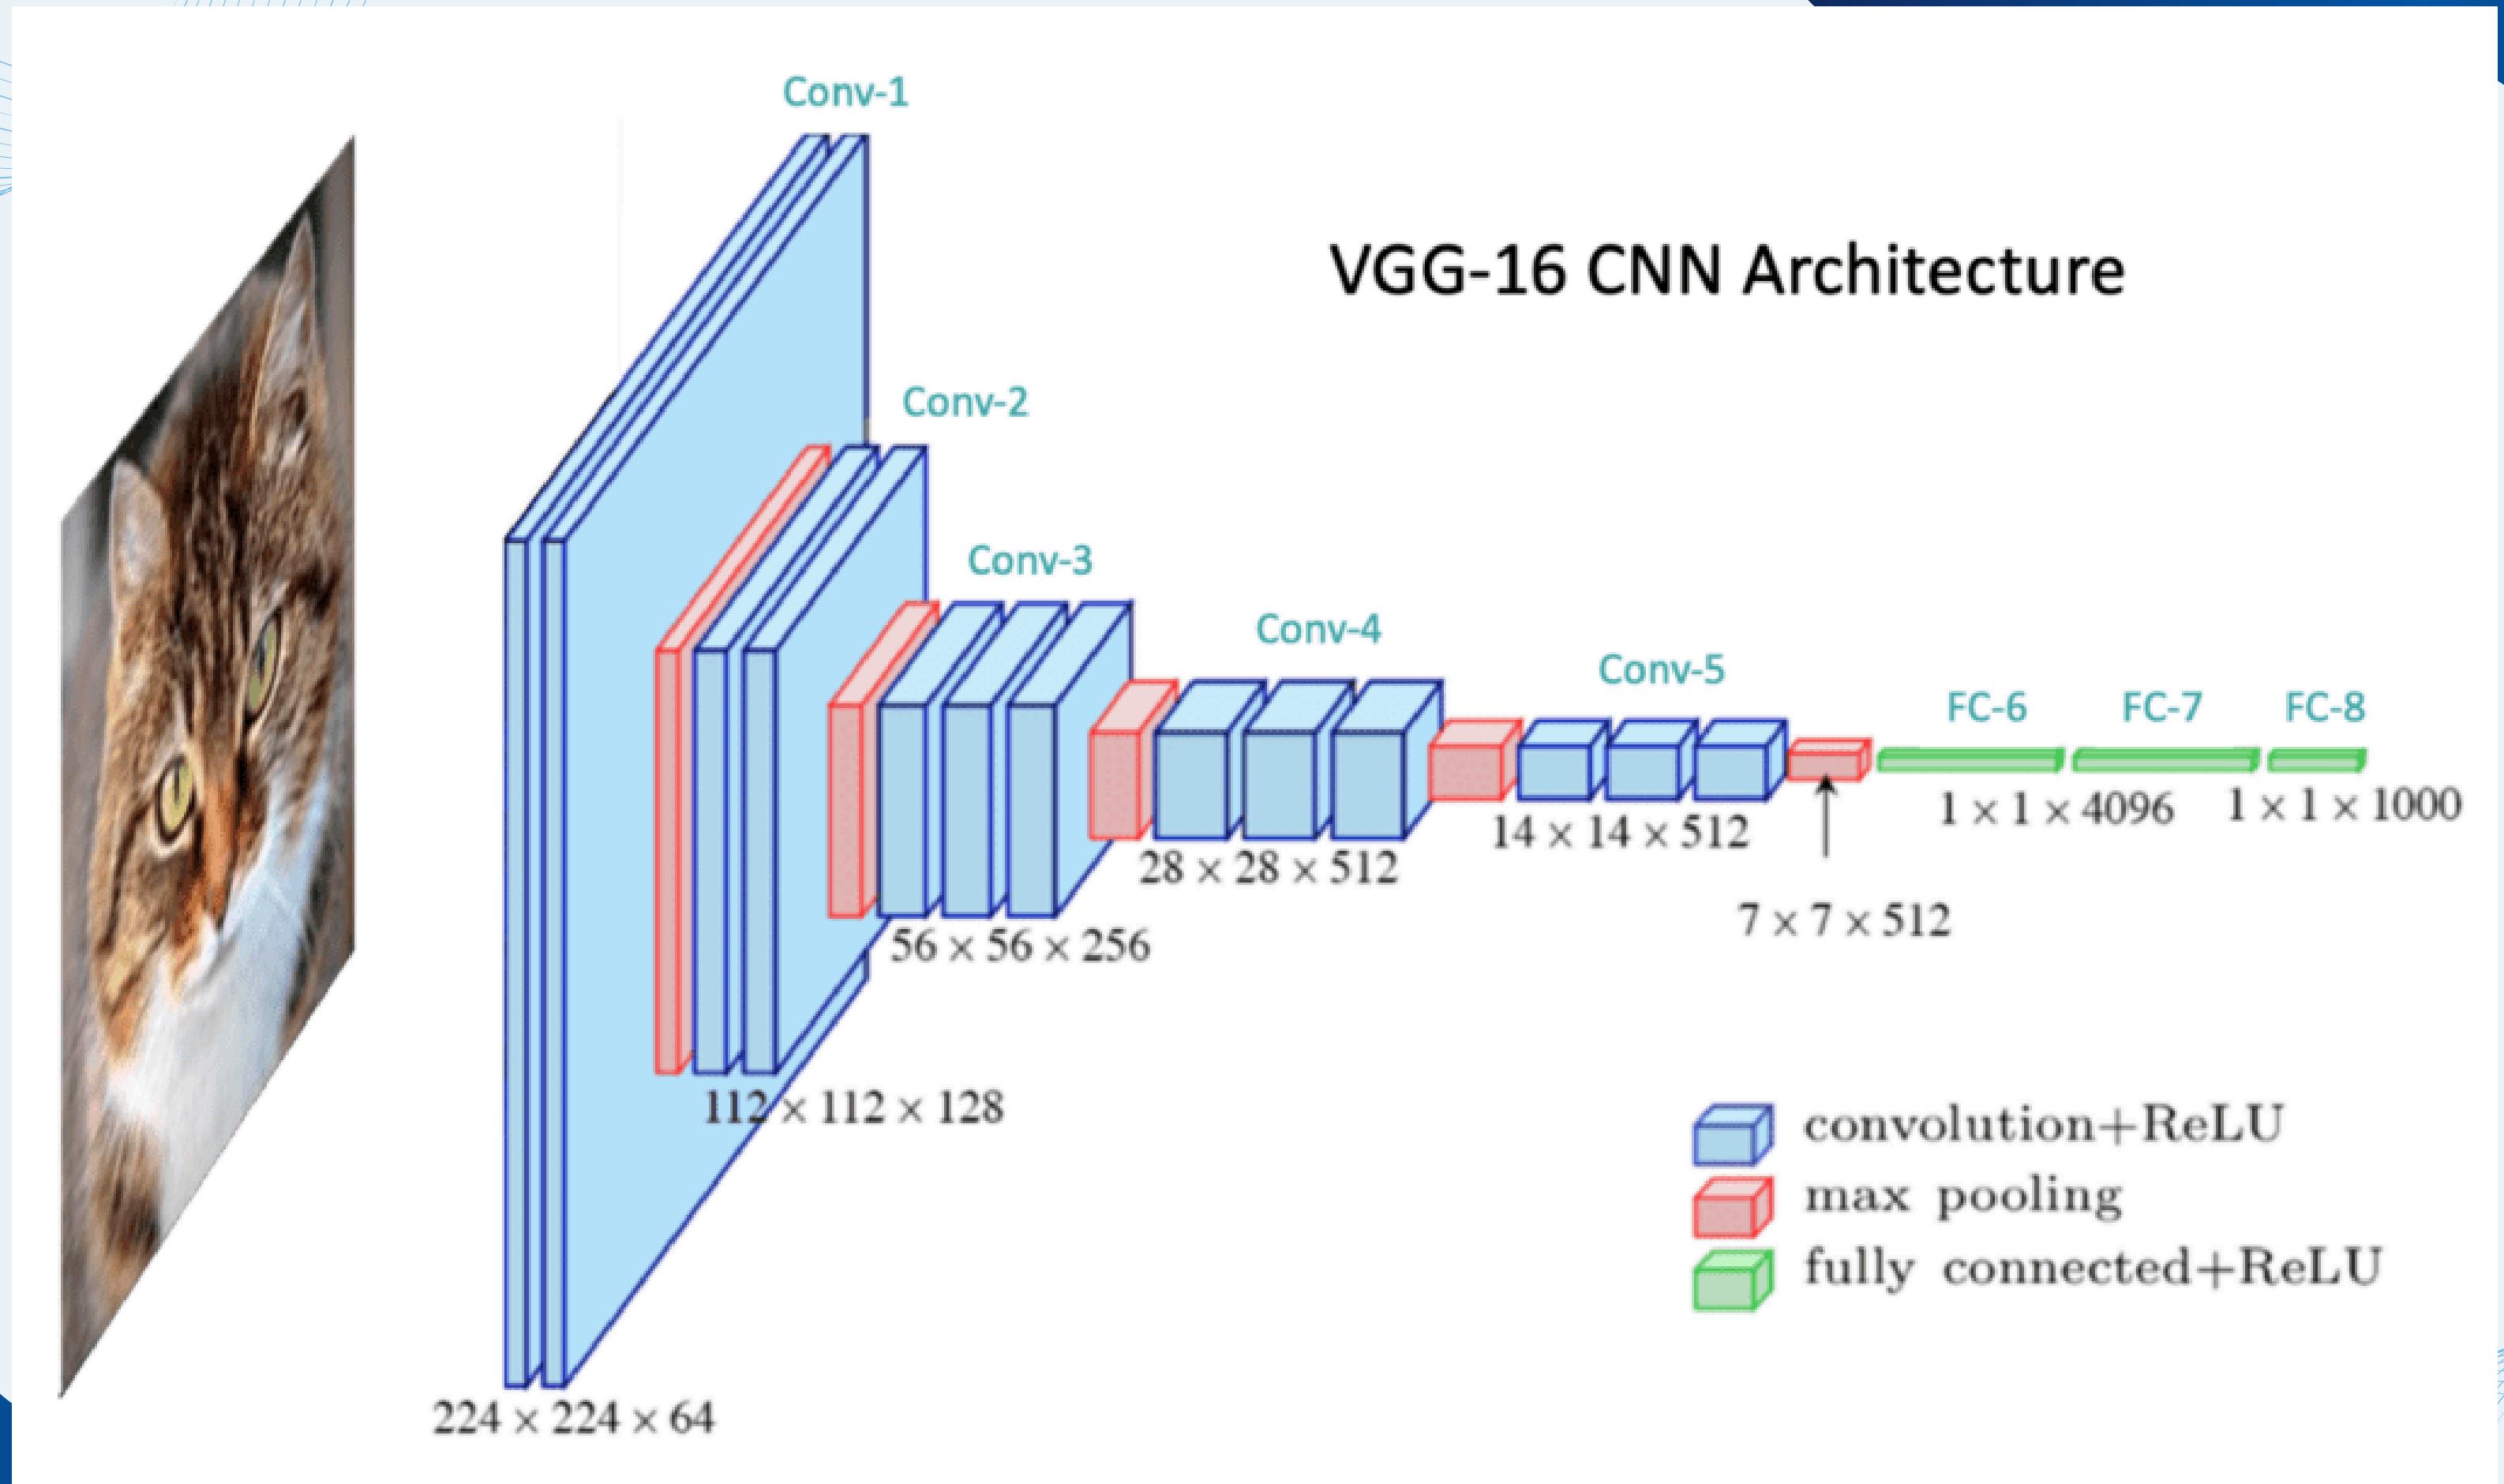
RETINAL DISEASE

=

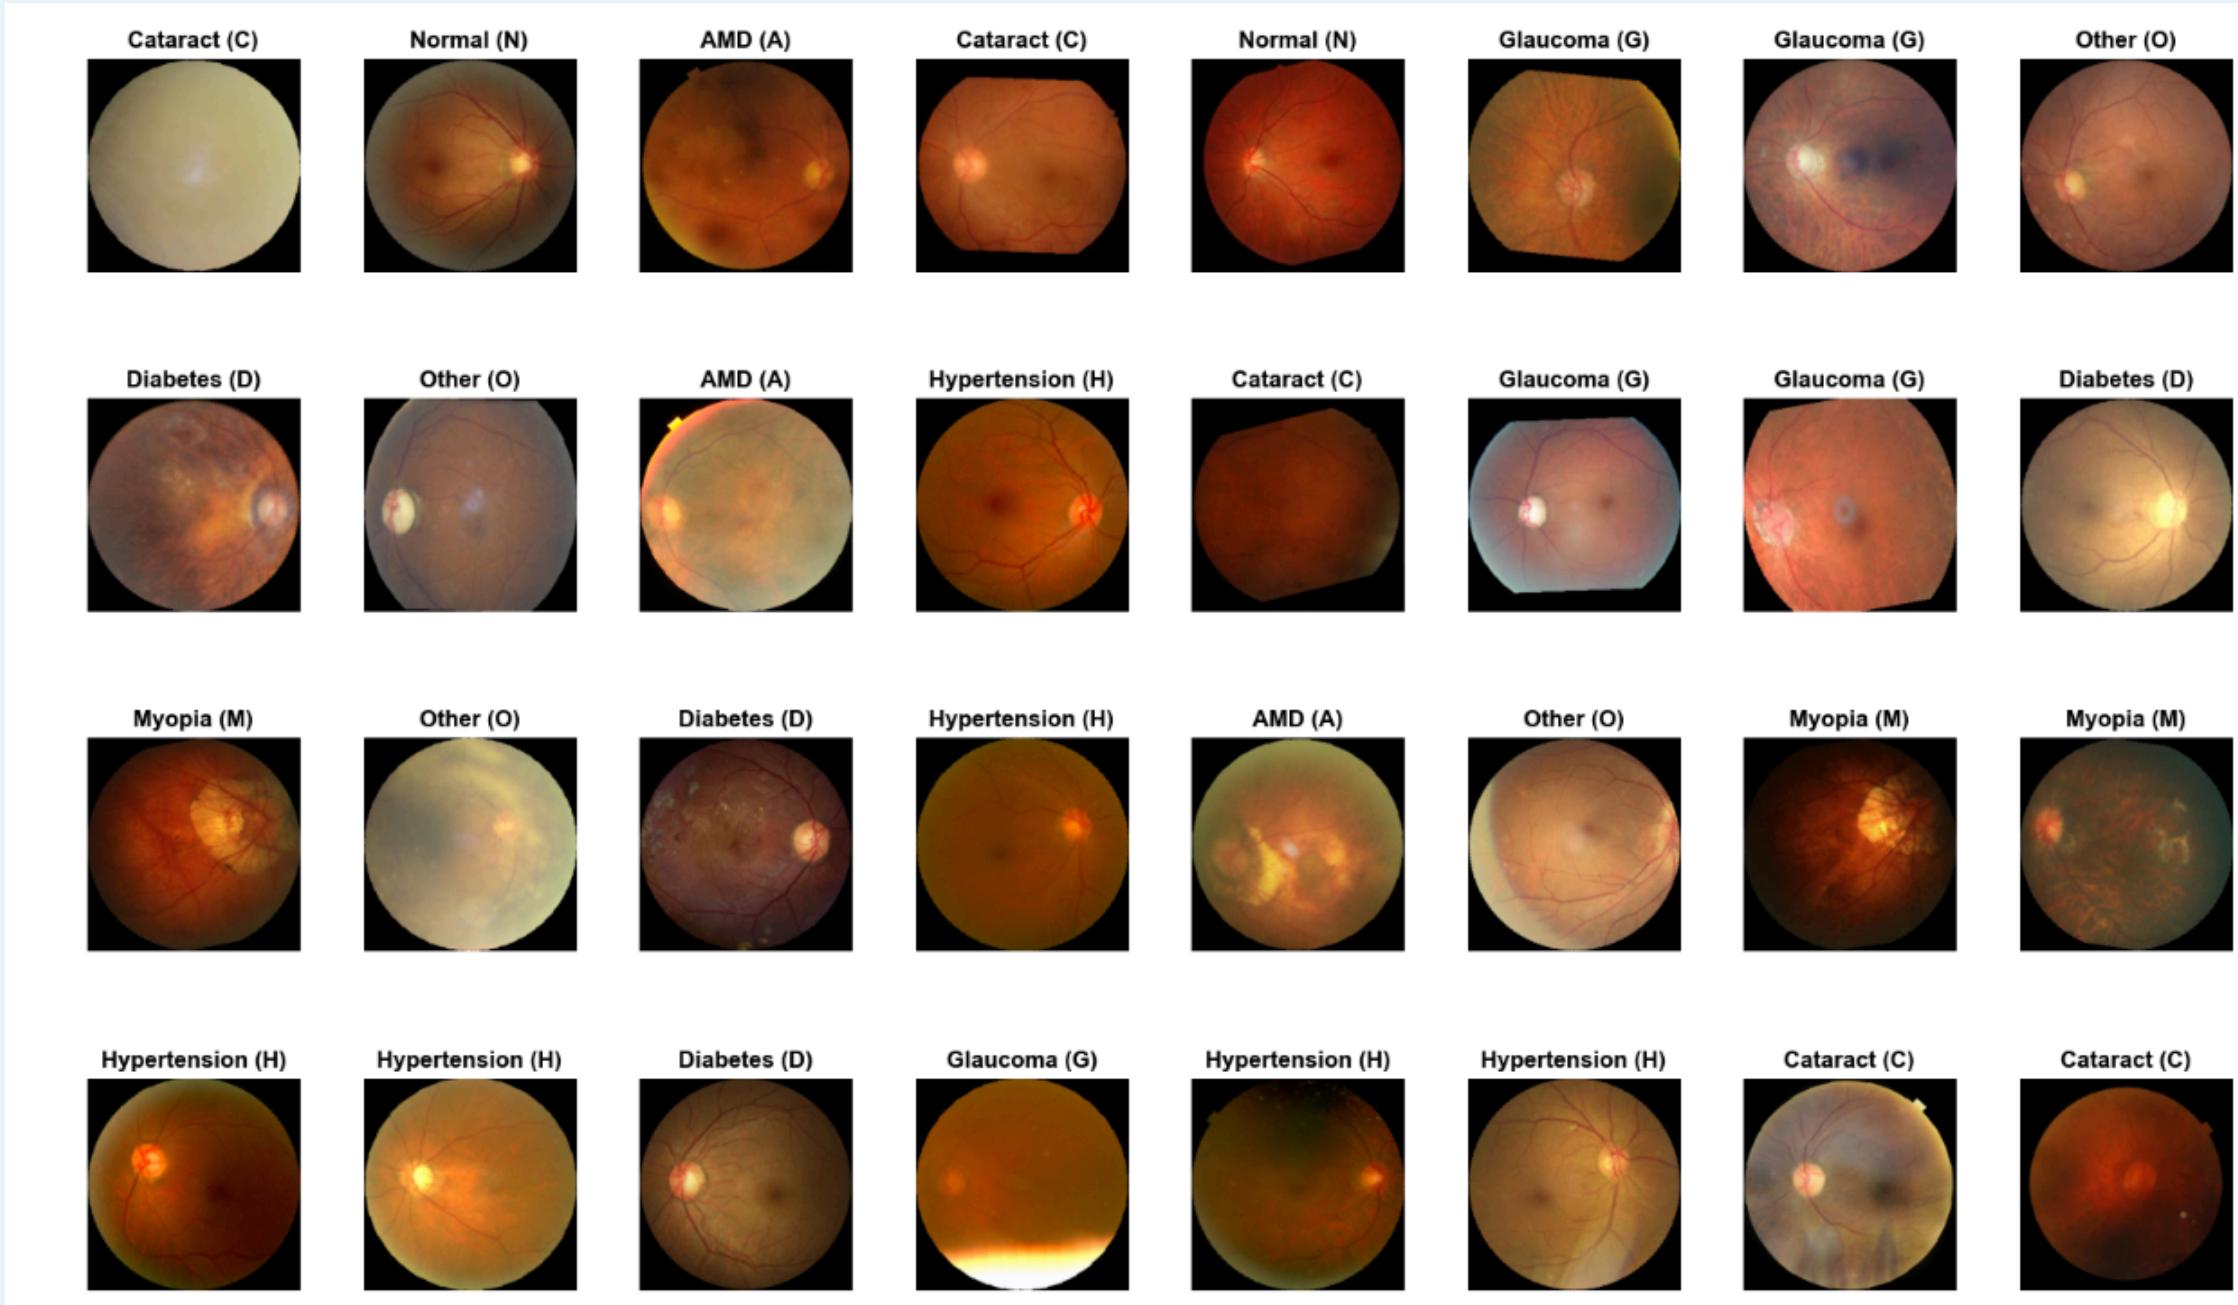
**LEADING CAUSES OF PREVENTABLE BLINDNESS
GLOBALLY.**



Complicated!



IDENTIFY INPUTS AND OUTPUTS



~10000 images of
retinal fundus images

For efficiency:
 $512 \times 512 \rightarrow 254 \times 254$

O: 'Normal (N)', 1: 'Diabetes (D)', 2: 'Glaucoma (G)', 3: 'Cataract (C)', 4: 'AMD (A)', 5: 'Hypertension (H)', 6: 'Myopia (M)', 7: 'Other (O)'

IDENTIFY INPUTS AND OUTPUTS

file csv

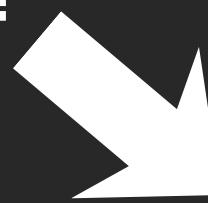
raw_df = pd.read_csv('label_images.csv')
raw_df.head()
✓ 0.1s

Python

	ID	Patient Age	Patient Sex	Left-Fundus	Right-Fundus	Left-Diagnostic Keywords	Right-Diagnostic Keywords	N	D	G	C	A	H	M	O	filepath	label	target	images
0	0.0	69.0	Female	0_left.jpg	0_right.jpg	cataract	normal fundus	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0/input/ocular-disease-recognition-odir5k/ODI...	[N]	[1, 0, 0, 0, 0, 0, 0]	0_right.jpg
1	1.0	57.0	Male	1_left.jpg	1_right.jpg	normal fundus	normal fundus	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0/input/ocular-disease-recognition-odir5k/ODI...	[N]	[1, 0, 0, 0, 0, 0, 0]	1_right.jpg
2	2.0	42.0	Male	2_left.jpg	2_right.jpg	laser spot, moderate non proliferative retinopathy	moderate non proliferative retinopathy	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0/input/ocular-disease-recognition-odir5k/ODI...	[D]	[0, 1, 0, 0, 0, 0, 0]	2_right.jpg
3	4.0	53.0	Male	4_left.jpg	4_right.jpg	macular epiretinal membrane	mild nonproliferative retinopathy	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0/input/ocular-disease-recognition-odir5k/ODI...	[D]	[0, 1, 0, 0, 0, 0, 0]	4_right.jpg
4	5.0	50.0	Female	5_left.jpg	5_right.jpg	moderate non proliferative retinopathy	moderate non proliferative retinopathy	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0/input/ocular-disease-recognition-odir5k/ODI...	[D]	[0, 1, 0, 0, 0, 0, 0]	5_right.jpg

```
describe_df = raw_df[['ID', 'label', 'images']]  
describe_df.describe(include='all')
```

Missing Entries!



	ID	label	images
--	----	-------	--------

count	6392.000000	9868	9868
-------	-------------	------	------

unique	NaN	8	9868
--------	-----	---	------

top	NaN	[N] NEWOTHER_EXTRA_DATA993.png	
-----	-----	--------------------------------	--

freq	NaN	3230	1
------	-----	------	---

```
# Train-Test-Val Split
train_df, temp_df = train_test_split(raw_df, test_size=0.25, random_state=42, stratify=raw_df['label'])
val_df, test_df = train_test_split(temp_df, test_size=0.5, random_state=42, stratify=temp_df['label'])

print(f"Train size: {len(train_df)} (Used for training)")
print(f"Val size: {len(val_df)} (Used for validating)")
print(f"Test size: {len(test_df)} (Used for final score)")
```

```
Train size: 7401 (Used for training)
Val size: 1233 (Used for early stopping)
Test size: 1234 (Used for final score)
```

STRATIFIED SPLITTING: ENSURE FAIR DISTRIBUTION

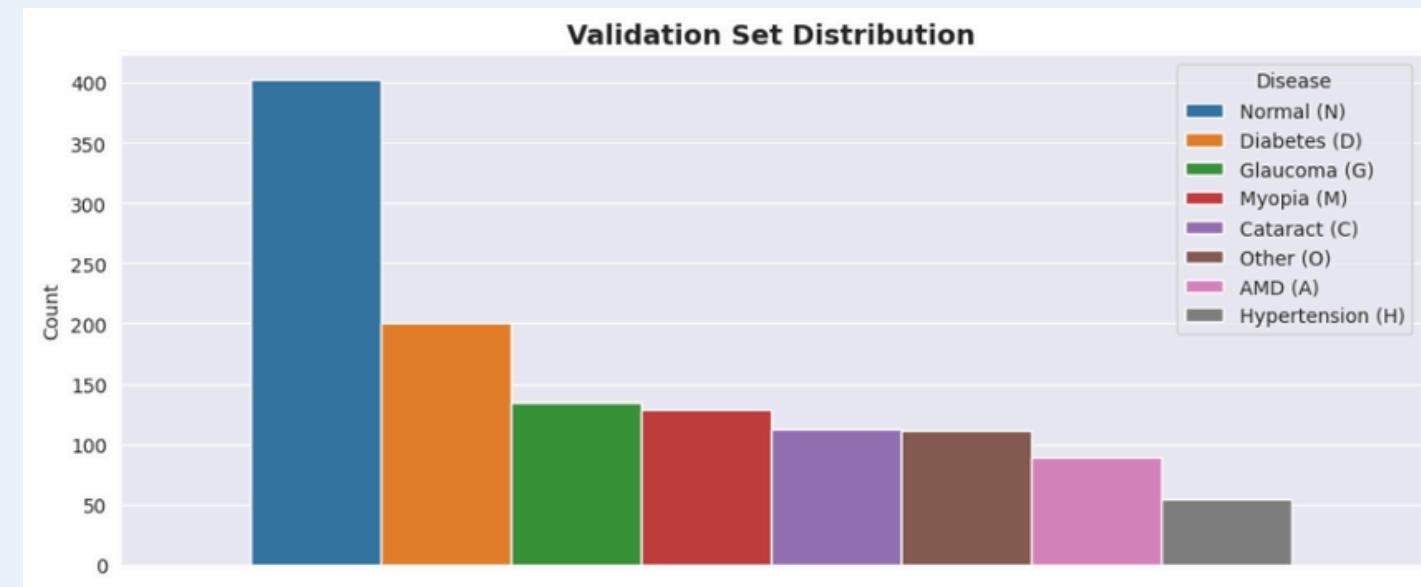
Training Set Distribution



Test Set Distribution



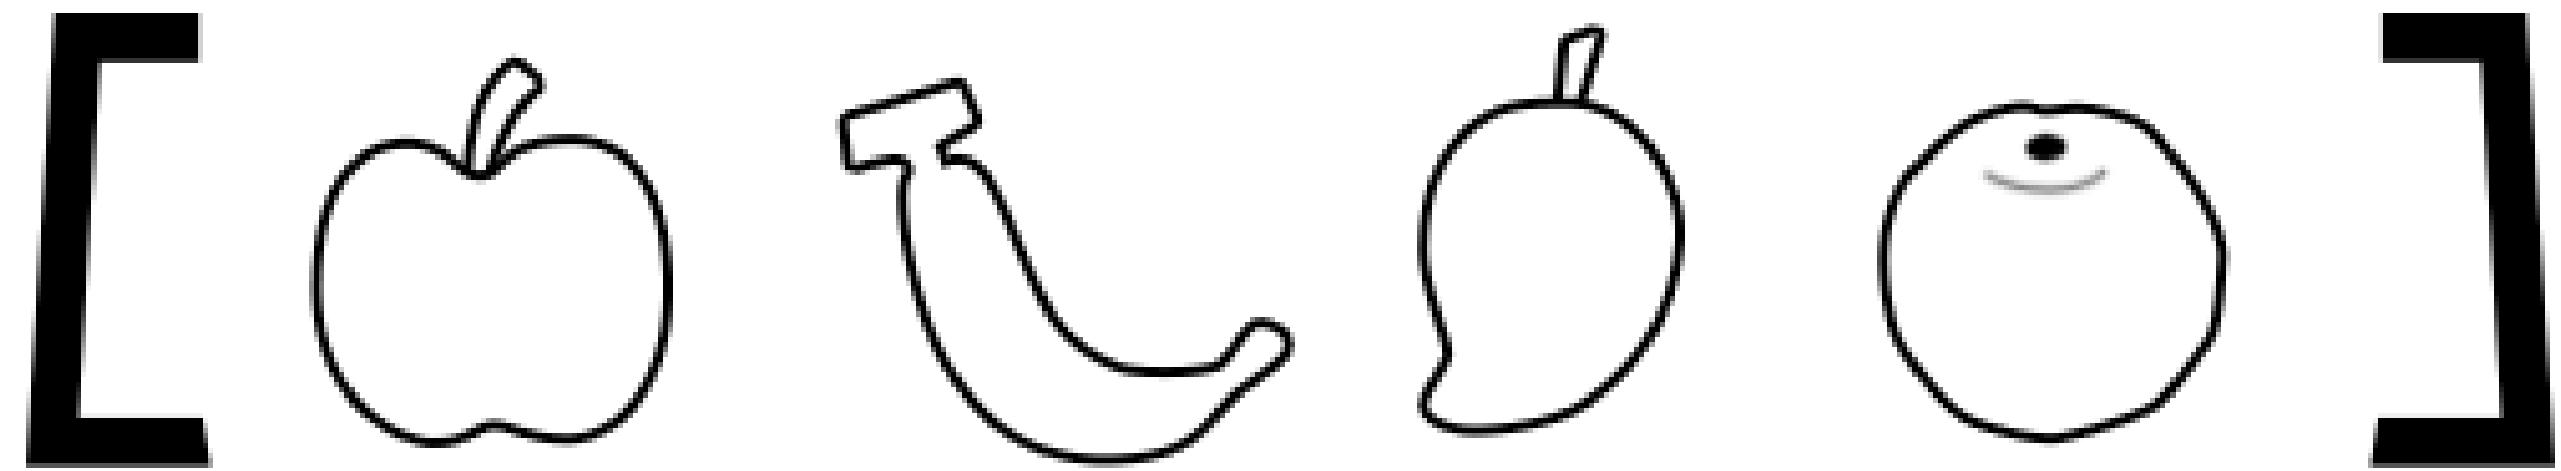
Validation Set Distribution



OPTIMIZATION TECHNIQUES

**WEIGHTED
RANDOM SAMPLER
(DATA LEVEL)**

**CLASS-WEIGHTED
LOSS FUNCTION
(ALGORITHM LEVEL)**



RARE

10%

30%

40%

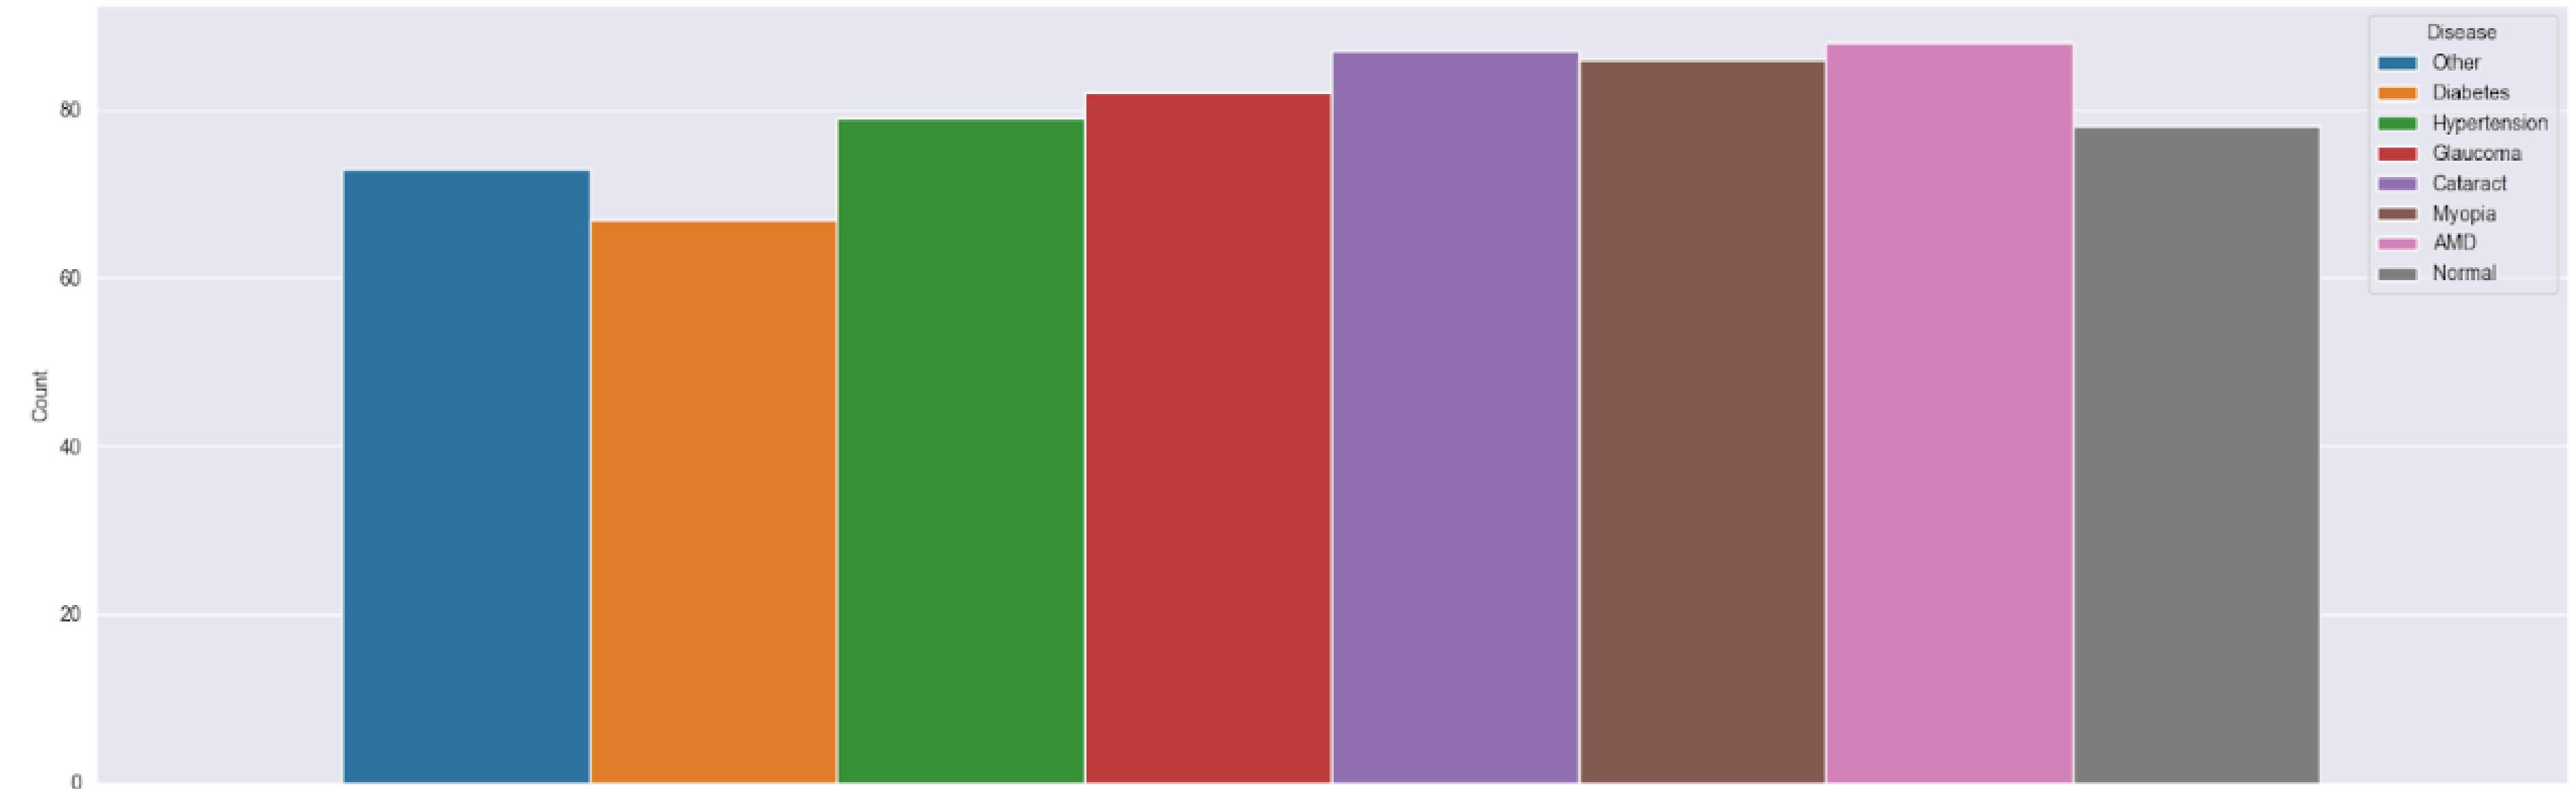
20%

HUGE WEIGHT

TINY WEIGHT

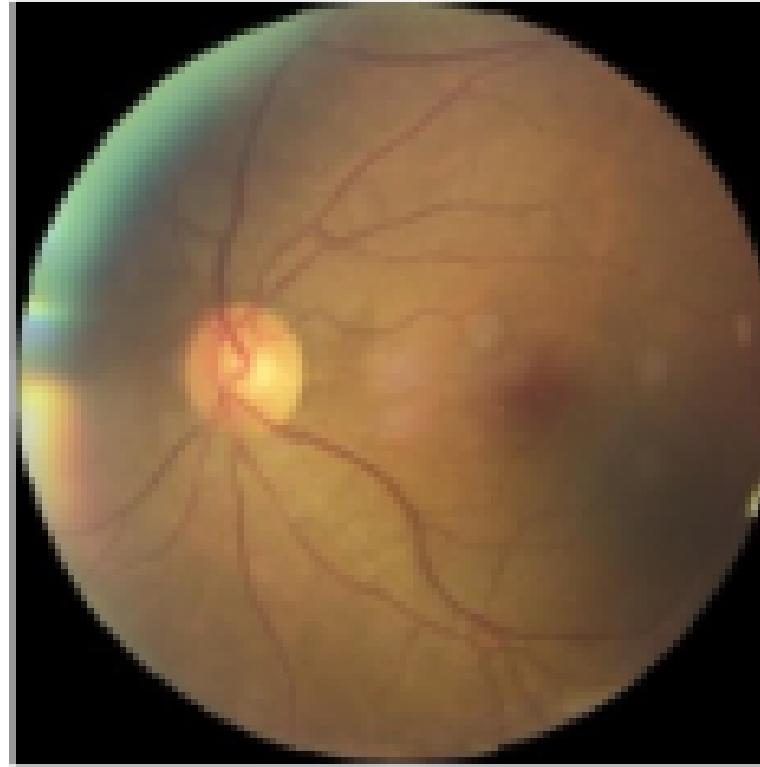
WEIGHTED RANDOM SAMPLER RESULT

Weighted Training Batches

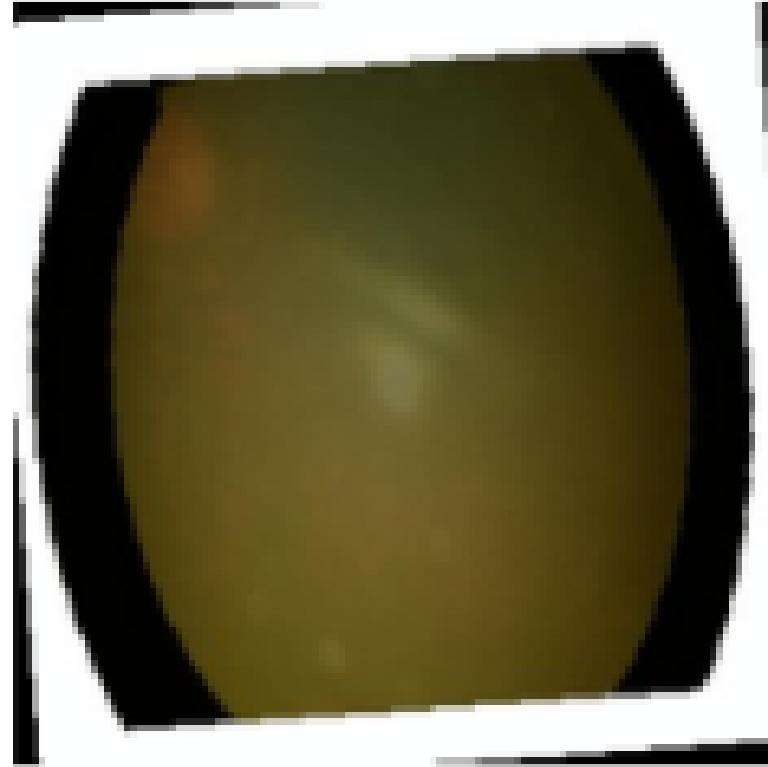


DATA AUGMENTATION

Other (0)



Cataract (C)



```
train_transforms = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(degrees=10),
    transforms.RandomAdjustSharpness(sharpness_t
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])
```

↑
ROTATED

NEURAL NETWORK ARCHITECTURE

```
def __init__(self, num_class):
    super(ClassificationCnn, self).__init__()
    self.num_class = num_class

    # 4 blocks to detect patterns.
    # Each reduces image size by half, doubles the number of "filters".
    self.conv = nn.Sequential([
        # Block 1
        # Input: (3, 256, 256) -> Output: (16, 128, 128)
        nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, padding=1),
        nn.BatchNorm2d(16),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),

        # Block 2
        # Input: (16, 128, 128) -> Output: (32, 64, 64)
        nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, padding=1),
        nn.BatchNorm2d(32),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),

        # Block 3
        # Input: (32, 64, 64) -> Output: (64, 32, 32)
        nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),

        # Block 4
        # Input: (64, 32, 32) -> Output: (128, 16, 16)
        nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),
    ])
```

NEURAL NETWORK ARCHITECTURE

```
# Takes the features found above and makes a decision.  
# Input Size: 128 channels * 6 width * 6 height  
self.fc = nn.Sequential(  
    # Layer 1: Connection from feature maps to hidden neurons  
    nn.Linear(128 * 16 * 16, 256),  
    nn.ReLU(),  
  
    # Dropout: Prevents the model from overfitting  
    nn.Dropout(0.25),  
  
    # Layer 2: Compress info further  
    nn.Linear(256, 128),  
    nn.ReLU(),  
  
    # Output: Final decision scores for each class  
    nn.Linear(128, self.num_class)  
)
```

NEURAL NETWORK ARCHITECTURE

```
# Takes the features found above and makes a decision.  
# Input Size: 128 channels * 6 width * 6 height  
self.fc = nn.Sequential(  
    # Layer 1: Connection from feature maps to hidden neurons  
    nn.Linear(128 * 16 * 16, 256),  
    nn.ReLU(),  
  
    # Dropout: Prevents the model from overfitting  
    nn.Dropout(0.25),  
  
    # Layer 2: Compress info further  
    nn.Linear(256, 128),  
    nn.ReLU(),  
  
    # Output: Final decision scores for each class  
    nn.Linear(128, self.num_class)  
)
```

NEURAL NETWORK ARCHITECTURE

```
def forward(self, x):
    """
    Forward pass defines how data moves through the graph.
    """

    # Extract features
    x = self.conv(x)

    # Flatten: Turn (Batch, 128, 6, 6) -> (Batch, 128)
    x = x.view(x.size(0), -1)

    # Classify
    x = self.fc(x)

    return x
```



TRAINING

```
# Criterion: Measures how wrong the model is
criterion = nn.CrossEntropyLoss()

# Optimizer: Adjusts the weights based on the errors.
optimizer = torch.optim.SGD(model.parameters(), lr=1e-2, momentum=0.9)

# Scheduler: Lower learning_rate if no improvement
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.1, patience=3)
```

TRAINING

```
# Training

# Track corrects and loss per epoch to calculate averages
train_corrects = 0
train_loss = 0.0

model.train()

# Wrap with tqdm for progress bar
for inputs, targets in tqdm(train_loader, desc=f'Training Epoch: {epoch + 1}/{N_EPOCHS}'):

    inputs, targets = inputs.to(device), targets.to(device)

    # Reset gradients for each batch
    optimizer.zero_grad()

    # Forward Pass: Ask model for prediction
    outputs = model(inputs)
```

TRAINING

```
# Validation

# Disable gradient calculation & dropout, lock batch norm
with torch.no_grad():
    val_corrects = 0
    val_loss = 0.0
    model.eval()
```

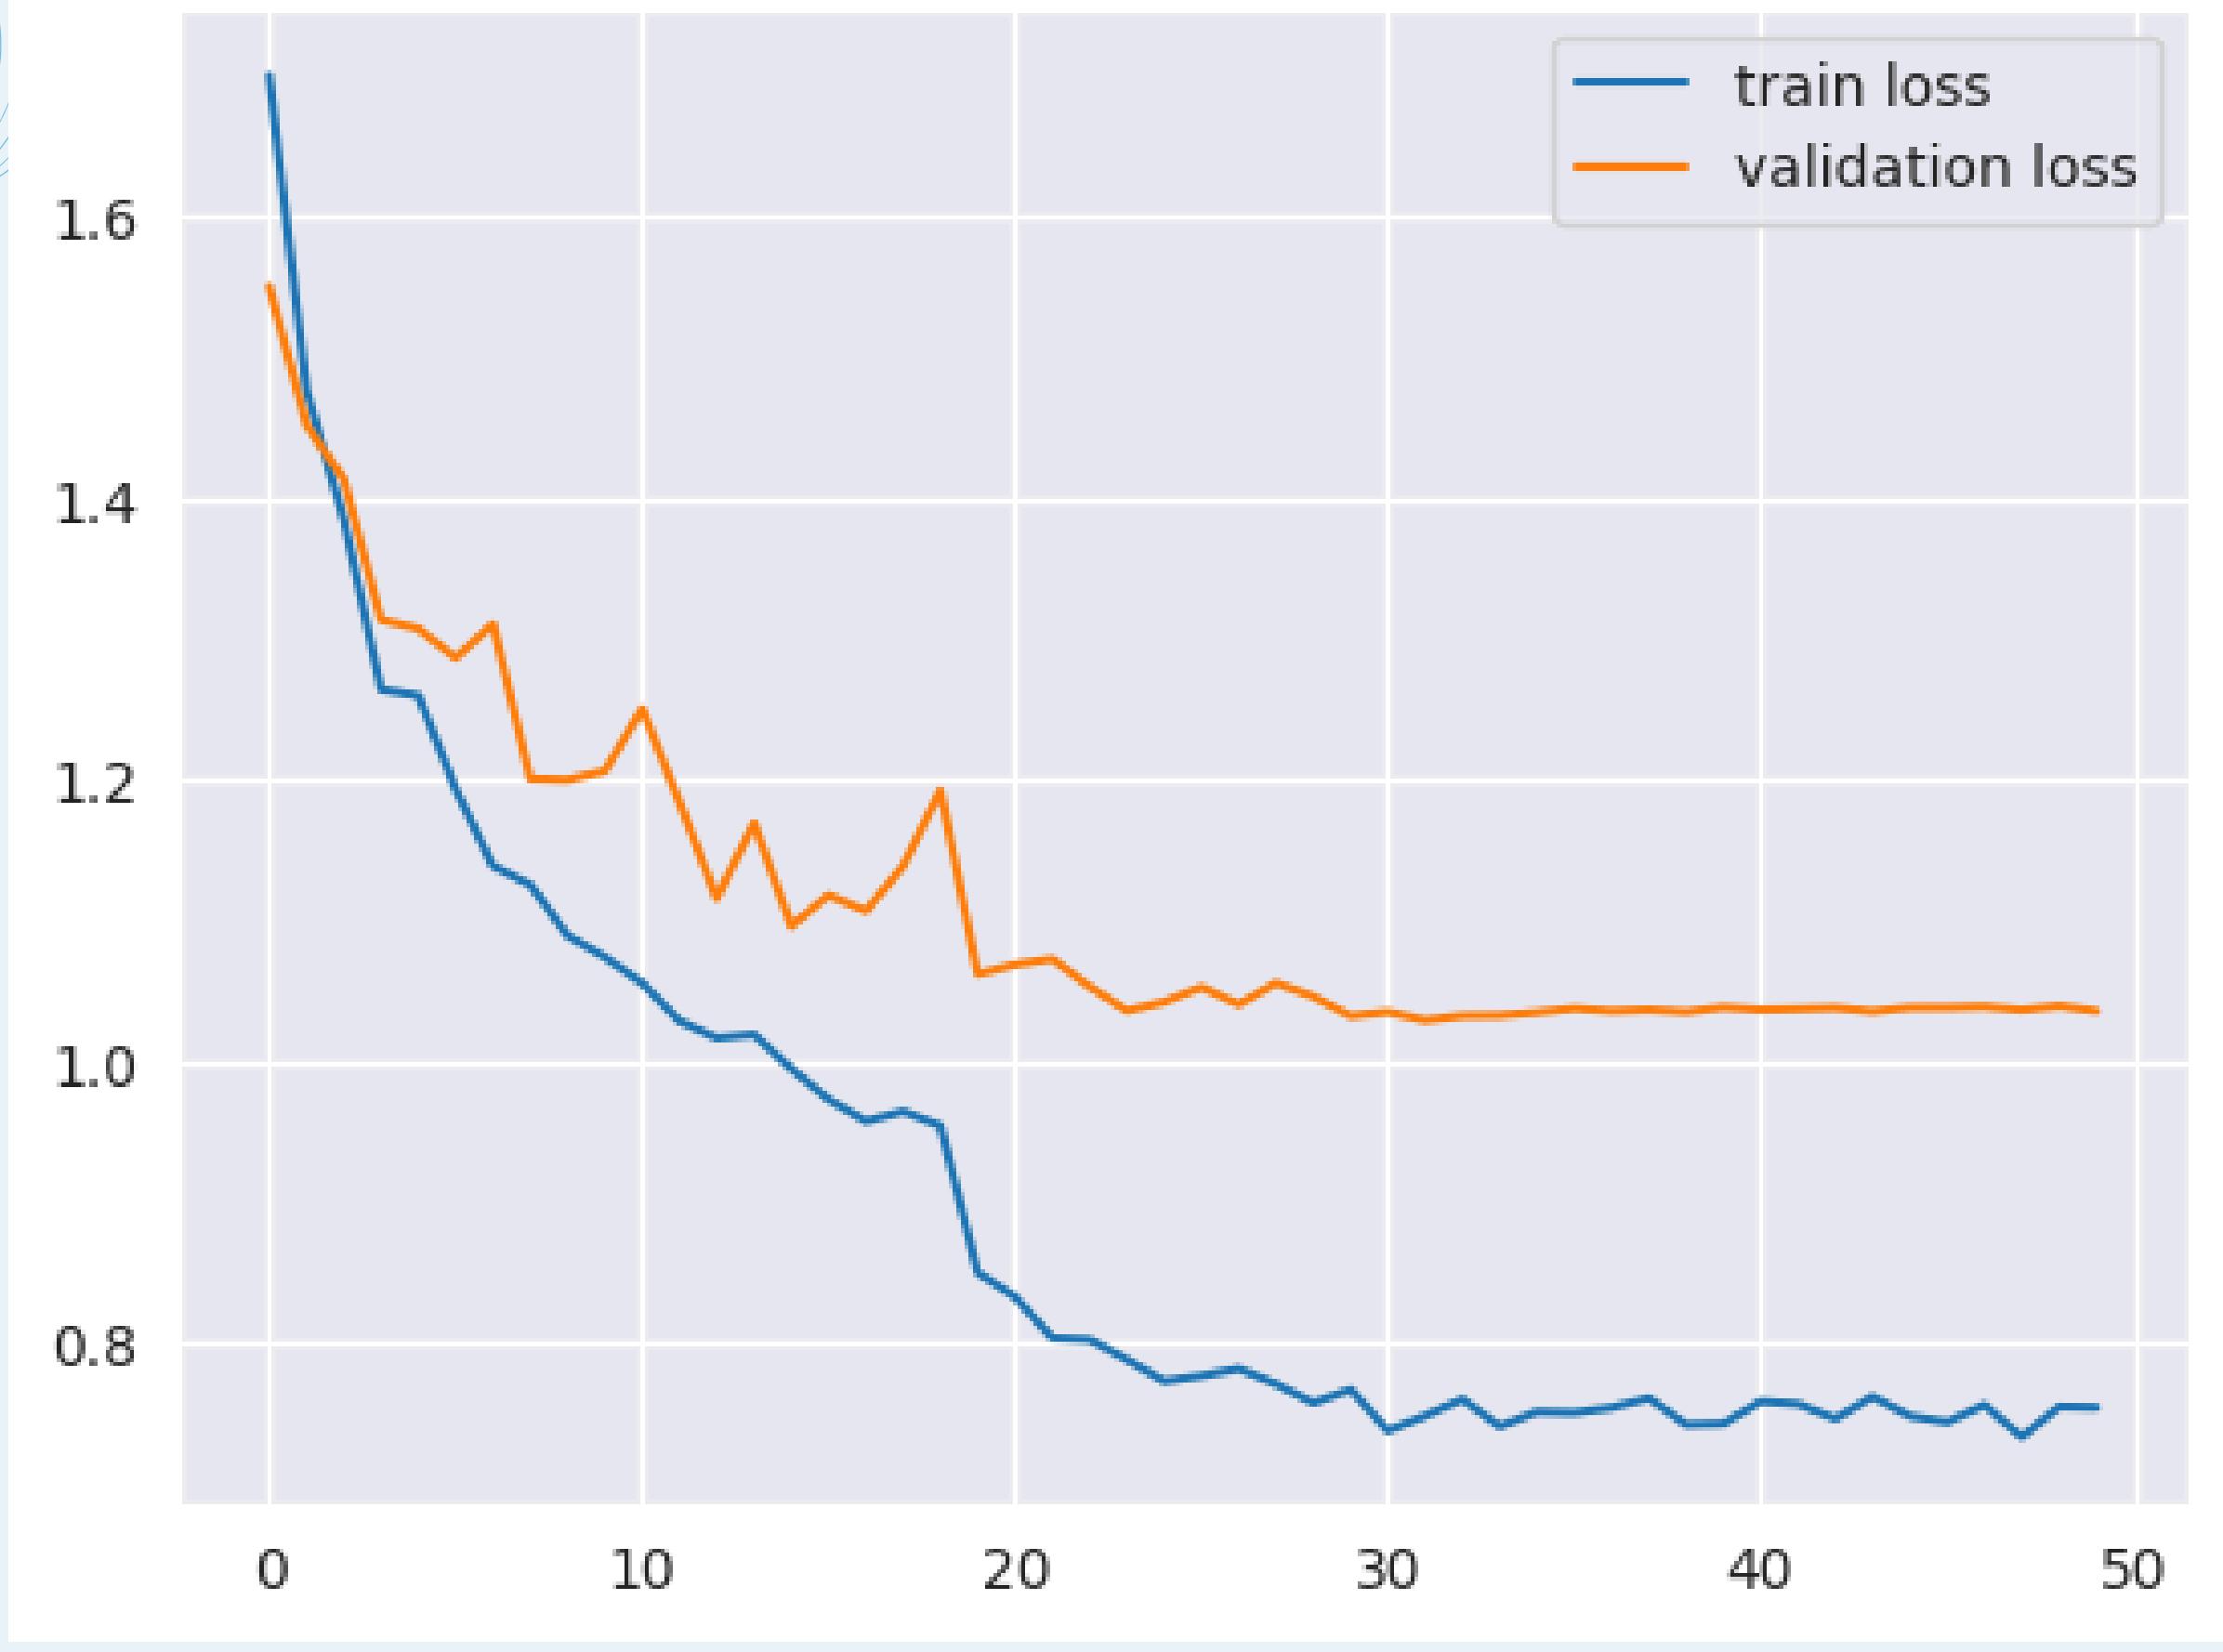
```
    val_loss = val_loss / len(val_loader.dataset)
    val_acc = val_corrects / len(val_loader.dataset)

    scheduler.step(val_loss)

    if val_acc > best_val_acc:
        best_val_acc = val_acc

    # Save copy of performing weights
    best_model_wts = copy.deepcopy(model.state_dict())
```

EVALUATION



Train Accuracy: 0.5426, Validation Accuracy: 0.5556, Test Accuracy: 0.5478

```
[[281 14 38 43 6 8 5 9]
 [145 17 2 12 1 12 2 10]
 [ 24 0 103 4 0 2 1 0]
 [ 2 0 4 88 16 0 2 0]
 [ 18 2 5 10 43 2 7 1]
 [ 11 0 3 1 1 38 0 0]
 [ 4 1 8 2 11 1 103 0]
 [ 63 5 1 15 20 2 2 3]]]
```

	precision	recall	f1-score	support
0	0.51	0.70	0.59	404
1	0.44	0.08	0.14	201
2	0.63	0.77	0.69	134
3	0.50	0.79	0.61	112
4	0.44	0.49	0.46	88
5	0.58	0.70	0.64	54
6	0.84	0.79	0.82	130
7	0.13	0.03	0.04	111
accuracy			0.55	1234
macro avg	0.51	0.54	0.50	1234
weighted avg	0.51	0.55	0.50	1234

STRENGTH

- DATA PIPELINE
- STRATIFIED SPLITTING
- WEIGHTED RANDOM SAMPLING
- CUSTOM VGG-LIKE ARCHITECTURE FROM SCRATCH

WEAKNESSES

- RESOLUTION CONSTRAINTS
- THE 'OTHER' CLASS

FUTURE WORK

- TRANSFER LEARNING
- HIGHER RESOLUTION

THANK YOU